DATA LAKES WITH SPARK

BIG DATA: You need a distributed system of several servers to be able to work with the data.

Numbers everyone should know

FASTEST

CPU. "Brain of computer" Operation: 0.4ns

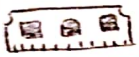

Memory/RAM. "ephemeral storage". Memory Reference: 100ns

Storage: SSD/Magnetic Disk. Random Read: 16 µs

SLOWEST

Network: Access to outside. Round trip for data from EU to US: 150ms

2.5 GHz CPU: 2.5 Billion Operations / second.

if 1 operation = 8 bytes → CPU can process 20 Billion Bytes/sec.

Memory is 2x faster than storage

But it's more expensive.

Distributed Computing vs. Parallel Computing

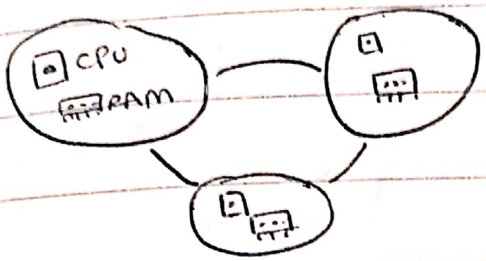

Several Nodes accross Network.
Each CPU has it's own memory

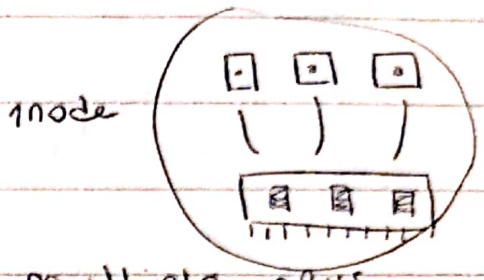

multiple CPUs share a single memory.

→ 200x Faster than RAM
→ 15x " than SSD
→ 20x " than network

Hadoop Framework

HDFS - Data Storage
MapReduce - Data Processing
YARN - Resources Manager
Hadoop Commons - libraries + Utilities

(+) other TOOLS on TOP.

Pig - SQL for MR
Hive - SQL for MR
SPARK
Storm - StreamigData
Flink - " "

MR

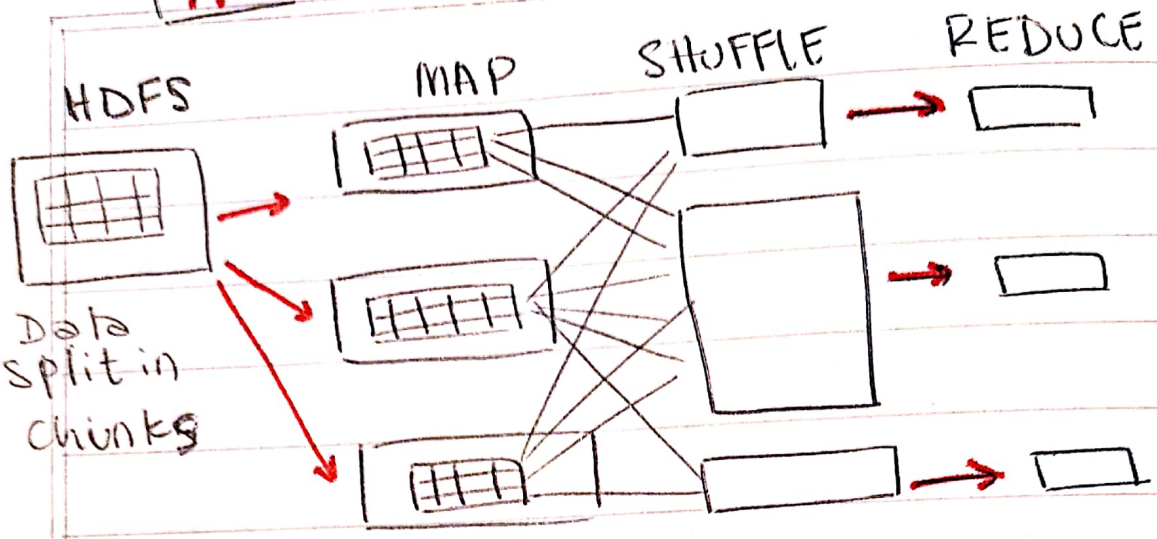

A B B
A C B B
C D A
D E

key, values → keys, value.

(A, 1) (B, 1) (B, 1) (A, 1) (C, 1) ...

(A, 1) (A, 1) (A, 1) (B, 1) (B, 1) (B, 1) (B, 1) ...

(A, 3) (B, 4) (C, 2) (D, 2) (E, 1)

DAG: directed Acyclical Graph

Spark builds a step by step directions of what functions + data it will need

→ Once it builds the DAG from the code it checks if there is anything it can procrastinate

→ LAZY EVALUATION

IMMUTABLE

Spark does not change of mutate data

MAPS

Makes a copy of the original input data and applies a function to it.

example

```
import pyspark
sc = pyspark.SparkContext(appName="...")
distrib_songs = sc.parallelize(songs)
distrib_songs.map(lambda x: x.lower())

distrib_songs.collect()
```

SPARK PROGRAMS

- entry point

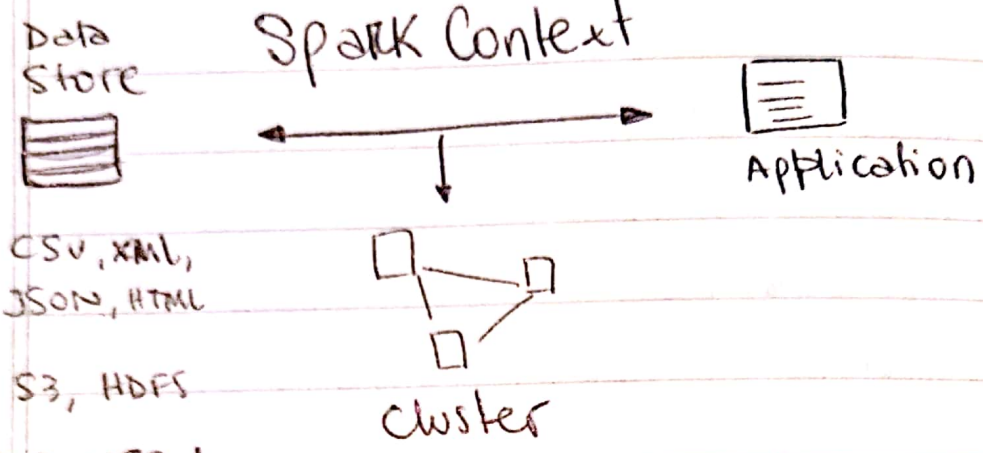


```
from pyspark import SparkContext, SConf
config = SparkConfig().setAppName(" ")
    .setMaster(" IP llocal")
SC = SparkContext(conf = conf)
```

TO read DF:

SC ↗ SQL equiv.

```
from pyspark.sql import SparkSession
Spark = SparkSession.builder /
        .appName(" ").config(" ")./
        .getOrCreate()
```

Data wrangling with DF

Read JSON into DF:

```
user_log = spark.read.json(path)
user_log.printSchema()
"        .describe()
         .show(n=1)
         .take(5)
```

SAVE into CSV:

```
"   .write.save(out_path, format="csv",
                header=True)
user_2 = spark.read.csv(path, header=True)
       .describe('column').show()
       .count()          .dropDuplicates()
       .select()         .sort("column")
       .where(column == value)
       .groupby( )
```

draw a Plot of your df

```
pd = df.toPandas()
plt.scatter(pd[x], pd[y])
plt.xlim(-1, 24);
func = udf(lambda x: ..., IntegerType())
df.withColumn("colName", func(col))
```

DATA LAKES

- Abundance of unstruct. data
- Unprecedented data volumes
- Rise of Big Data techs.
- Data analysis - new types (ex. ML, NLP,)
- Emergence of new roles (ex. Data Scientist)

*Evolved from Data warehouse to cope with:
- Variety of data formats and structuring
- Support to agile + ad-hoc data exploration.
- Wider data transformation needed

Big Data tech:

(i) ETL Offloading : instead of ETL grid or staging area → same Hw for storage and processing.

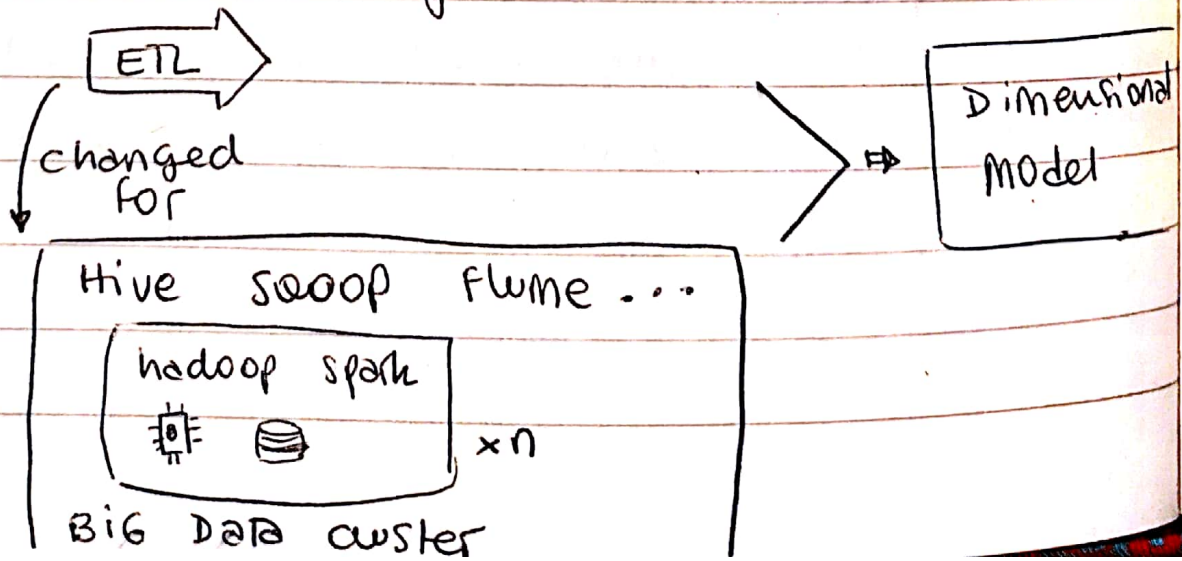

(11) schema-on-read → As easy to work with files as with a database, without having to → create a DB
→ insert data into DB.

```
df = spark.read.csv(path,
        inferSchema = "True", header=True,
        sep=';').
```

To specify a schema:

```
schema = StructType([
        StructField("col-name", IntegerType()),
        ... ])
df = spark.read.csv(path,
     schema=schema, sep=';',
     mode='DROPMALFORMED')

df.createOrReplaceTempView("Table")
spark.sql("select * from table").show(5)
```

(III) Un Structured support

- Spark can read/write files in
  - Text based formats → Text, JSON, CSV..
  - Binary formats → Avro, Parquet ..
  - Compressed formats → Gzip, Snappy

- Read/write from <> FS:
  - Local file system
  - HDFS
  - S3

- & Databases:
  - SQL through JDBC
  - NoSQL → MongoDB, Cassandra, Neo4j

Spark.read.format("jdbc"). ...

	Data Warehouse	Data Lake
Data Format	Tabular	All formats
Ingestion	ETL	ELT
Data model	Star snowflake w/ conformed dimensions or DM or OLAP cubes	star, snow, OLAP + ad-hoc reps.
Schema	schema-on-write	schema-on-read
Tech	Expensive MPP, disks, netws.	Commodity Hw.
Data Q.	High, consistent, clean	Mixed, raw, some transfo.
users	BA	DS, BA, ML engineers
Analysis	Reports, BI viz.	ML, Graphs, data exploration

DATA LAKE ON AWS

Storage	Processing	AWS-managed solutions	Vendor Managed

HDFS → Spark → AWS EMR (HDFS + spark, spark only)

S3 → Spark → EC2 + vendor sol.

S3 → Serverless → AWS Athena → serverless + vendor solution

AWS EMR (HDFS + Spark)

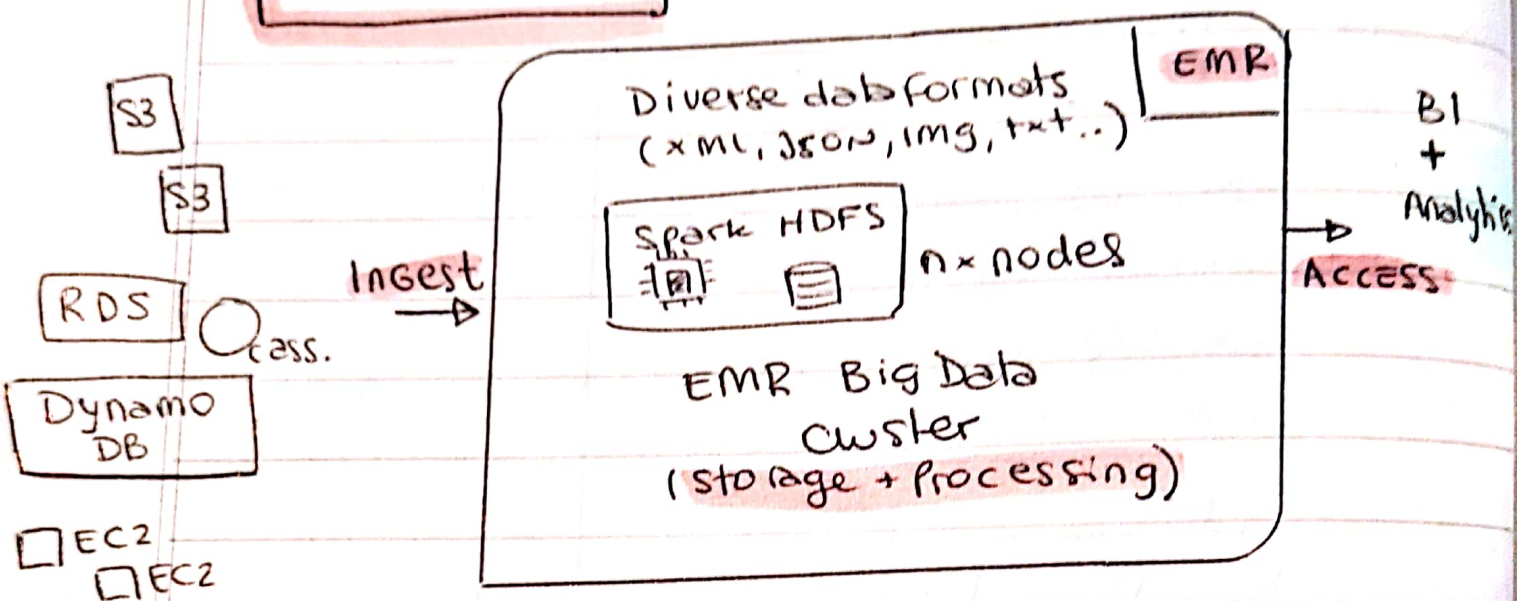

Sources

- Once ingested all data is stored on HDFS, processed on the cluster + accessed from Analytics + BI Apps
- EMR cluster is created once, it can grow but it cannot be shutdown or else it loses all the stored data.

! COST → keeping the cluster running -

(S3 + Spark)

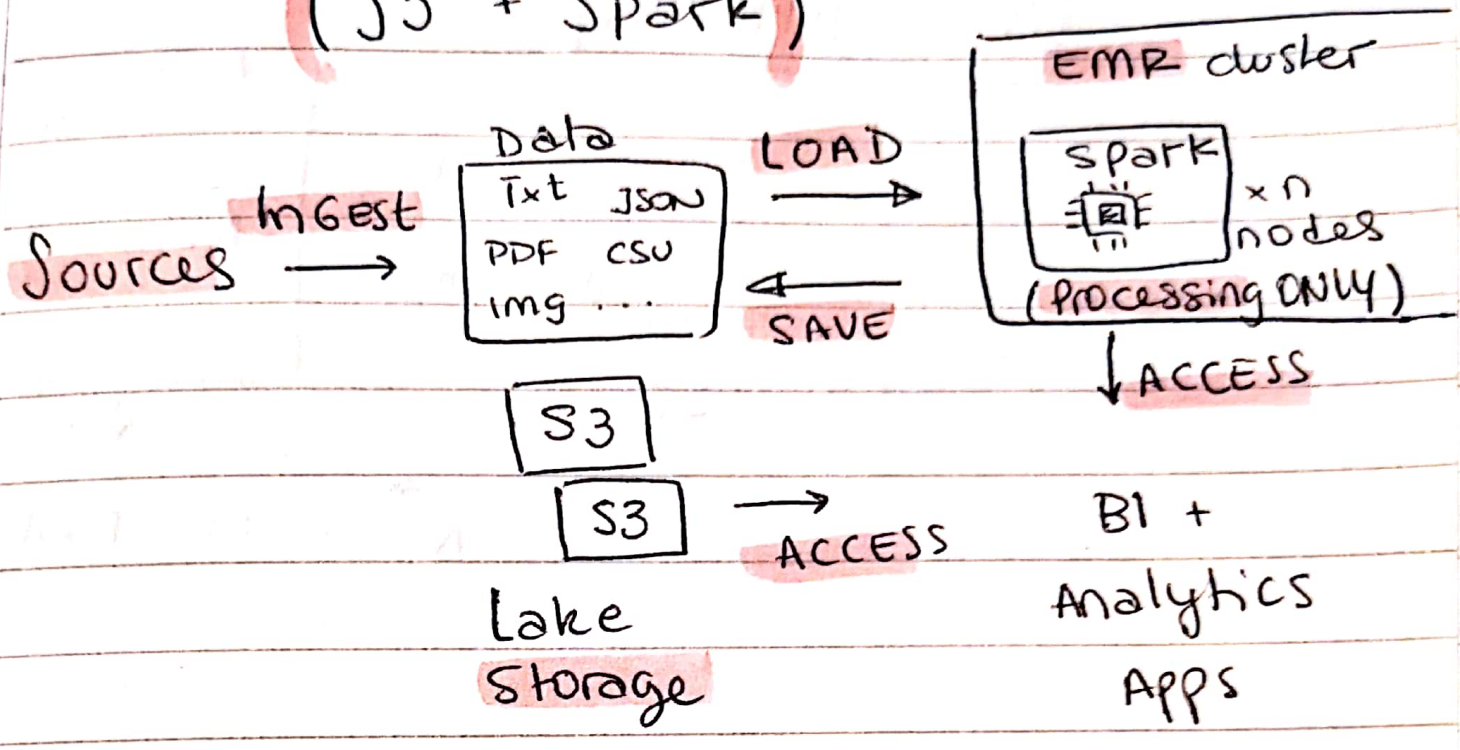

- NO HDFS in EMR
  - EMR cluster is spun on-demand, shutdown otherwise.

Potentially → less costly, easier to manage, more performant.

AWS ATHENA

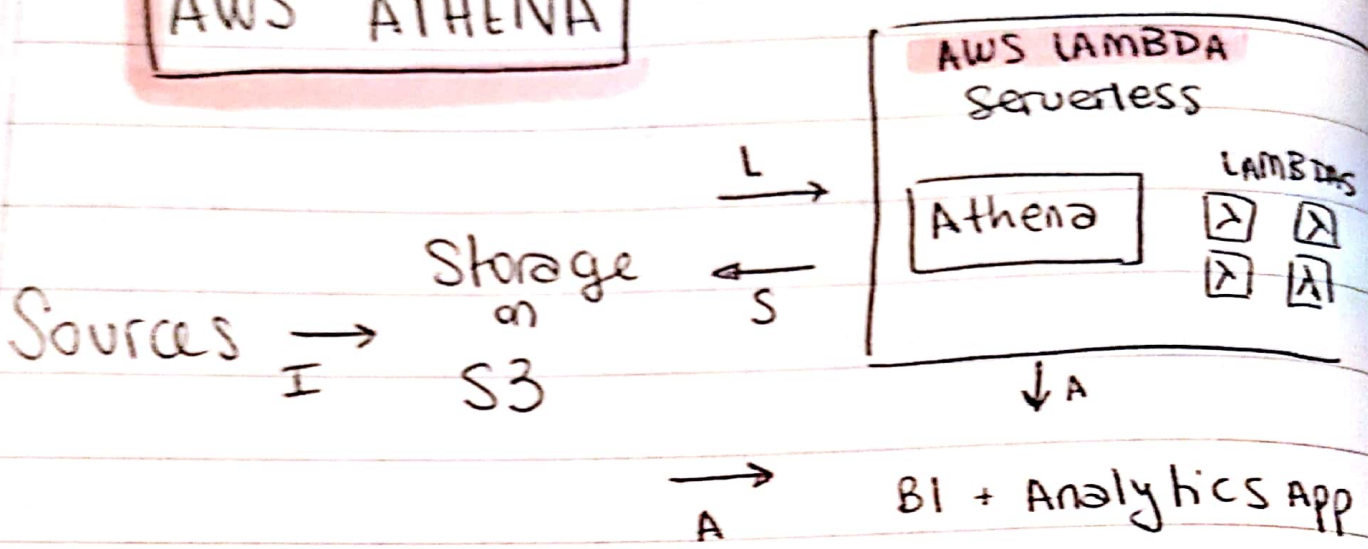

LAMBDA → Function as a Service
(Pay by execution time)

ATHENA → loads + processes data on serverless lambda resources

- Transparent management of resources